

Decode Time per Token (4 GPU Average)

